

LOAN APPROVAL ANALYSIS

```
In [163]: # Importing the important Libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns

#tensorflow Lib
import tensorflow
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
In [164]: # Loading the dataset
df = pd.read_json("C:/Users/Likhith Gaikwad/Downloads/loan_approval_c
df.head(10)
```

Out[164]:

		Id	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownership
0	1	1303834	23		3	single	rented	no
1	2	7574516	40		10	single	rented	no
2	3	3991815	66		4	married	rented	no
3	4	6256451	41		2	single	rented	yes
4	5	5768871	47		11	single	rented	no
5	6	6915937	64		0	single	rented	no
6	7	3954973	58		14	married	rented	no
7	8	1706172	33		2	single	rented	no
8	9	7566849	24		17	single	rented	yes
9	10	8964846	23		12	single	rented	no

Exploratory Data Analysis (EDA)

In [165]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 252000 entries, 0 to 251999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    252000 non-null int64
1   Income                252000 non-null int64
2   Age                  252000 non-null int64
3   Experience             252000 non-null int64
4   Married/Single        252000 non-null object
5   House_Ownership       252000 non-null object
6   Car_Ownership         252000 non-null object
7   Profession            252000 non-null object
8   CITY                  252000 non-null object
9   STATE                 252000 non-null object
10  CURRENT_JOB_YRS       252000 non-null int64
11  CURRENT_HOUSE_YRS    252000 non-null int64
12  Risk_Flag             252000 non-null int64
dtypes: int64(7), object(6)
memory usage: 26.9+ MB

```

In [166]: df.describe()

Out[166]:

		Id	Income	Age	Experience	CURRENT_JOB_YRS
count	252000.000000	252000.000000	252000.000000	252000.000000	252000.000000	252000.000000
mean	126000.500000	4.997117e+06	49.954071	10.084437	6.33	
std	72746.278255	2.878311e+06	17.063855	6.002590	3.64	
min	1.000000	1.031000e+04	21.000000	0.000000	0.00	
25%	63000.750000	2.503015e+06	35.000000	5.000000	3.00	
50%	126000.500000	5.000694e+06	50.000000	10.000000	6.00	
75%	189000.250000	7.477502e+06	65.000000	15.000000	9.00	
max	252000.000000	9.999938e+06	79.000000	20.000000	14.00	

In [167]: *# Finding Null values*
df.isnull().sum()

```

Out[167]: Id                    0
Income                  0
Age                    0
Experience              0
Married/Single         0
House_Ownership        0
Car_Ownership          0
Profession             0
CITY                   0
STATE                  0
CURRENT_JOB_YRS        0
CURRENT_HOUSE_YRS      0
Risk_Flag              0
dtype: int64

```

```
In [170]: # Checking for duplicates in the data
duplicates = df[df.duplicated()]
print(duplicates)
```

Empty DataFrame

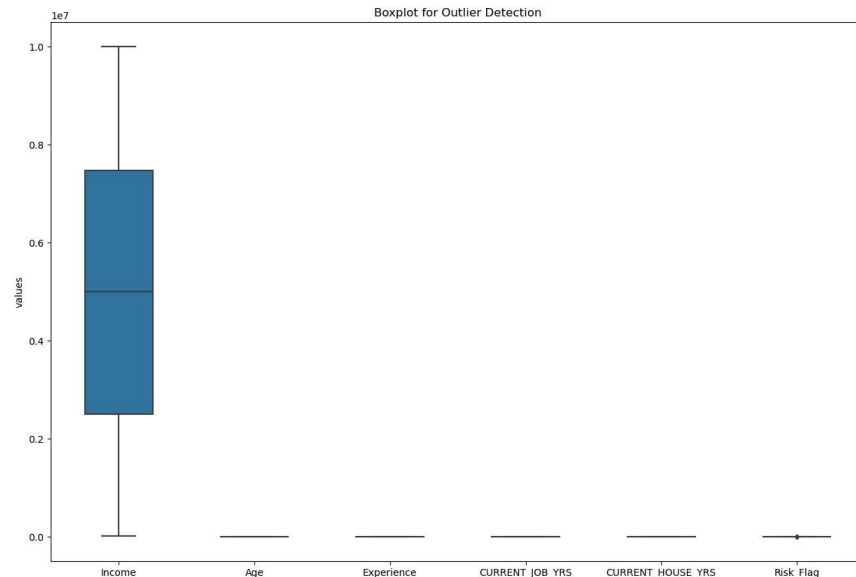
Columns: [Id, Income, Age, Experience, Married/Single, House_Ownership, Car_Ownership, Profession, CITY, STATE, CURRENT_JOB_YRS, CURRENT_HOUSE_YRS, Risk_Flag]

Index: []

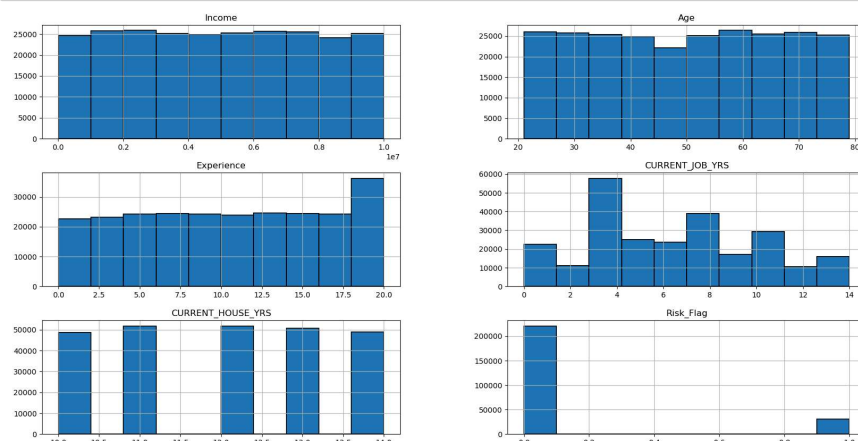
Data Visualization

```
In [171]: # Removing the ID column
df = df.drop(columns=['Id'])
```


```
In [172]: # Checking for outliers in the data
fig, ax = plt.subplots(figsize=(15,10))
sns.boxplot(data=df, width=0.5, ax=ax, fliersize=3)
ax.set_title('Boxplot for Outlier Detection')
ax.set_ylabel('values')
plt.show()
```



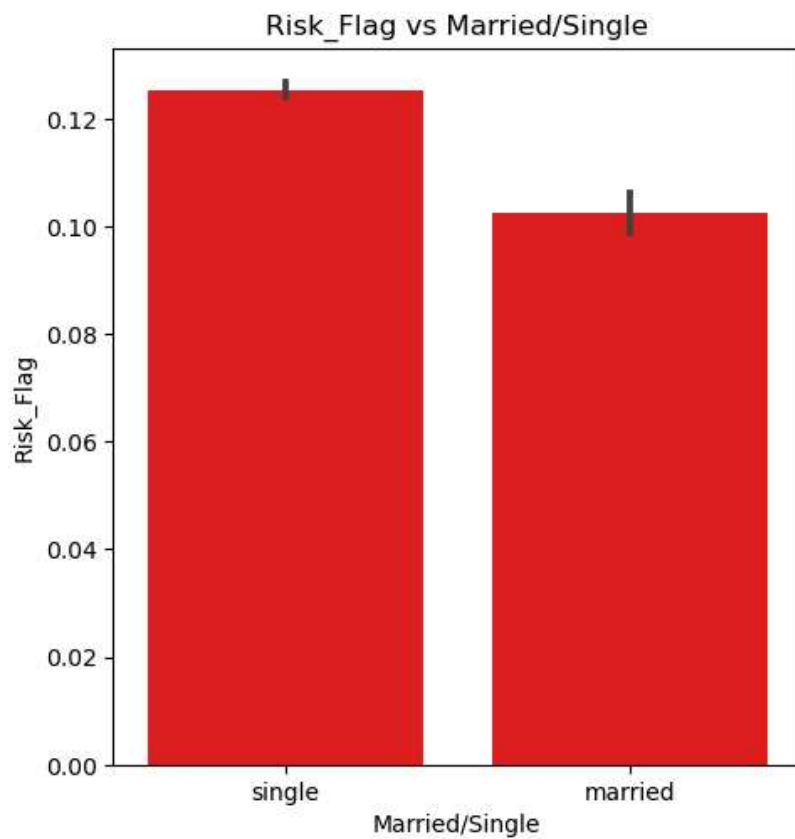
```
In [173]: # Histogram plot on different variables
df.hist(edgecolor='black', linewidth=1.2, figsize=(20,10))
plt.show()
```



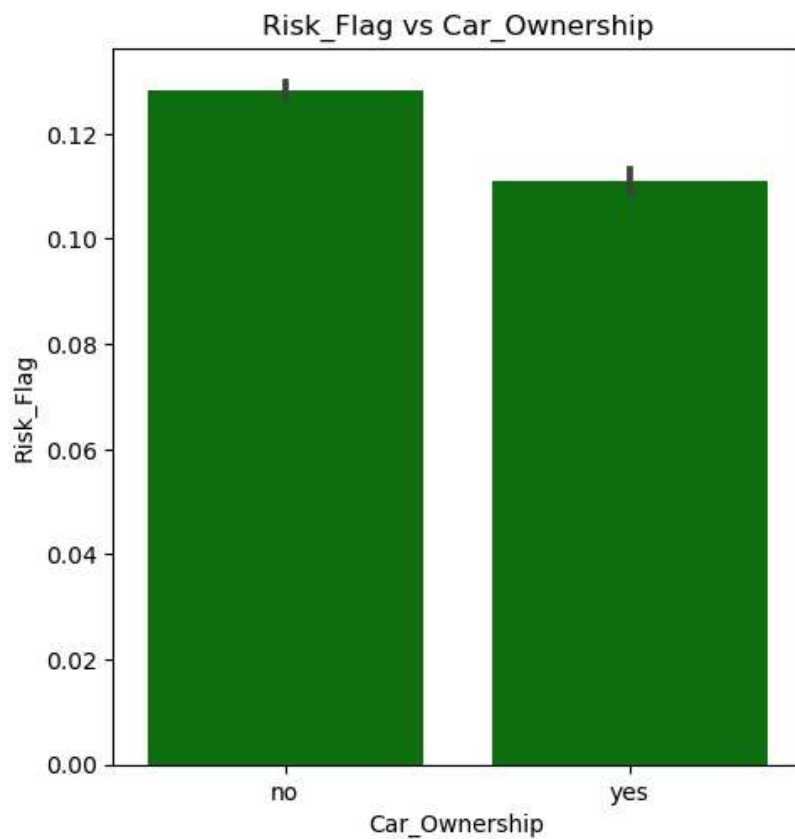
```
In [36]: # Barplot between Profession and Risk_Flag
plt.figure(figsize=(10,10))
sns.barplot(x= df['Profession'],y=df['Risk_Flag'], color='red')
plt.tight_layout()
plt.title("Risk_Flag vs Profession")
plt.xticks(rotation=90)
plt.show()
```



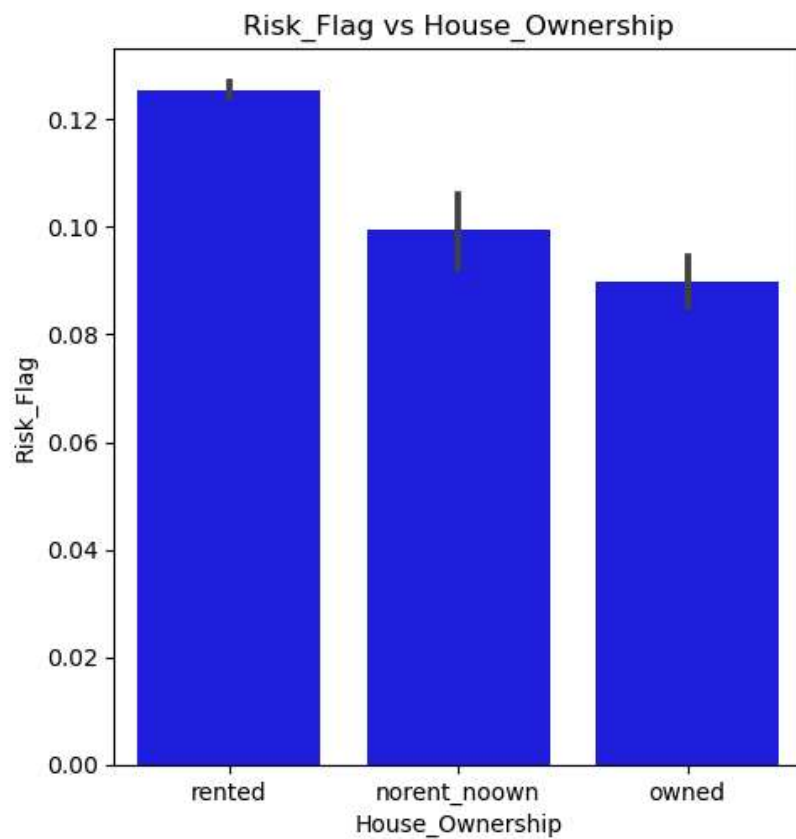
```
In [157]: # Barplot between Marital status and Risk_Flag
plt.figure(figsize=(5,5))
sns.barplot(x=df['Married/Single'],y=df['Risk_Flag'], color='red')
plt.tight_layout()
plt.title("Risk_Flag vs Married/Single")
plt.show()
```




```
In [158]: # Barplot between Car_Ownership and Risk_Flag
plt.figure(figsize=(5,5))
sns.barplot(x=df['Car_Ownership'],y=df['Risk_Flag'], color='green')
plt.tight_layout()
plt.title("Risk_Flag vs Car_Ownership")
plt.show()
```



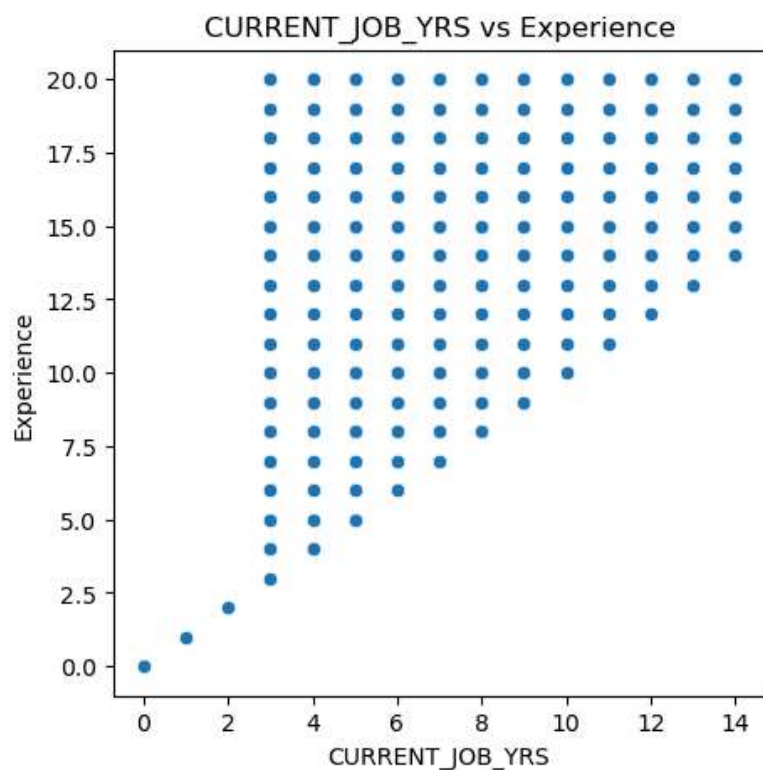
```
In [159]: # Barplot between House_Ownership and Risk_Flag
plt.figure(figsize=(5,5))
sns.barplot(x=df['House_Ownership'],y=df['Risk_Flag'], color='blue')
plt.tight_layout()
plt.title("Risk_Flag vs House_Ownership")
plt.show()
```




```
In [153]: # Identifying correlation between variables
corr = df.corr()
sns.heatmap(corr, annot=True, center=0, square=True, linewidth=0.5)
plt.show()
```



```
In [156]: # Scatterplot between CURRENT_JOB_YRS and Experience since they are
plt.figure(figsize=(5,5))
plt.title("CURRENT_JOB_YRS vs Experience")
sns.scatterplot(y='Experience', x='CURRENT_JOB_YRS', data=df)
plt.show()
```




```
In [154]: # Visualizing the target variable
sns.countplot(x='Risk_Flag', data=df)
plt.show()
```



```
In [40]: # There is a large imbalance in the target variable
# Making the target variable balanced

# Separating the classes
df_majority = df[df['Risk_Flag']==0]
df_minority = df[df['Risk_Flag']==1]

# Downsampling the majority class
df_majority_downsampled = df_majority.sample(len(df_minority), random_state=42)

# Combining the downsampled majority and minority classes
df_new = pd.concat([df_majority_downsampled, df_minority])

# Counting the Target variable
print(df_new['Risk_Flag'].value_counts())

0    30996
1    30996
Name: Risk_Flag, dtype: int64
```

```
In [82]: # Removing the target variable
X_old = df_new.drop(columns=['Risk_Flag'])
X_old
```

```
Out[82]:
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownershi
114032	8823581	27	15	single	rented	n
213746	5157518	23	6	single	rented	n
163964	4041403	47	19	single	rented	ye
155854	8622588	45	18	single	owned	n
101715	3803419	61	16	married	rented	n
...
251973	1244622	35	15	single	rented	n
251977	1330613	63	19	single	rented	n
251981	1796713	47	2	single	rented	n
251982	3182290	52	2	single	rented	n
251993	8141027	60	10	single	rented	n

61992 rows × 11 columns

```
In [83]: # Inverting the categorical variables to dummy/indicator variables
# Adding data to the X and y variables
pd.get_dummies(X_old, columns=['Married/Single', 'House_Ownership', 'Car_Ownership'])
df_new['Risk_Flag']
```

```
In [43]: # Splitting the train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [44]: # Scaling the data
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

```
In [45]: X_train.shape
```

```
Out[45]: (43394, 403)
```

Model Performance

```
In [190]: # Building the model

model = Sequential()

model.add(Dense(800, activation='relu', kernel_initializer='he_uniform'))
model.add(Dropout(0.2))

model.add(Dense(400, activation='sigmoid')) #hidden Layer 2
model.add(Dropout(0.5))

model.add(Dense(200, activation = 'relu', kernel_initializer='he_normal'))
model.add(Dropout(0.1))


model.add(Dense(100, activation = 'sigmoid')) #hidden Layer 4
model.add(Dropout(0.2))


model.add(Dense(1, activation='sigmoid')) # output Layer
```


```
In [191]: # Compiling the model


model.compile(optimizer='adam', loss='binary_crossentropy', metrics=
```


```
In [192]: # Training the data  
model.fit(X_train, y_train, epochs=30, batch_size=30)
```


Epoch 1/30
1447/1447  10s 6ms/step - accuracy: 0.5550 - loss: 0.6865


Epoch 2/30
1447/1447  8s 6ms/step - accuracy: 0.6958 - loss: 0.5823


Epoch 3/30
1447/1447  8s 5ms/step - accuracy: 0.7928 - loss: 0.4652


Epoch 4/30
1447/1447  8s 5ms/step - accuracy: 0.8285 - loss: 0.4132


Epoch 5/30
1447/1447  8s 6ms/step - accuracy: 0.8486 - loss: 0.3812


Epoch 6/30
1447/1447  8s 6ms/step - accuracy: 0.8597 - loss: 0.3611


Epoch 7/30
1447/1447  8s 6ms/step - accuracy: 0.8705 - loss: 0.3360


Epoch 8/30
1447/1447  8s 5ms/step - accuracy: 0.8787 - loss: 0.3234


Epoch 9/30
1447/1447  8s 5ms/step - accuracy: 0.8837 - loss: 0.3150


Epoch 10/30
1447/1447  8s 5ms/step - accuracy: 0.8863 - loss: 0.3004


Epoch 11/30
1447/1447  8s 5ms/step - accuracy: 0.8909 - loss: 0.2943


Epoch 12/30
1447/1447  8s 6ms/step - accuracy: 0.8953 - loss: 0.2843


Epoch 13/30
1447/1447  8s 6ms/step - accuracy: 0.8979 - loss: 0.2765


Epoch 14/30
1447/1447  8s 6ms/step - accuracy: 0.8997 - loss: 0.2728


Epoch 15/30
1447/1447  8s 6ms/step - accuracy: 0.9052 - loss: 0.2631


Epoch 16/30
1447/1447  8s 6ms/step - accuracy: 0.9057 - loss: 0.2643


Epoch 17/30
1447/1447  8s 6ms/step - accuracy: 0.9081 - loss: 0.2562


Epoch 18/30
1447/1447  8s 6ms/step - accuracy: 0.9103 - loss: 0.2456

Epoch 19/30
1447/1447  8s 6ms/step - accuracy: 0.9129 - loss: 0.2440

Epoch 20/30
1447/1447  8s 6ms/step - accuracy: 0.9118 - loss: 0.2449

Epoch 21/30
1447/1447  8s 6ms/step - accuracy: 0.9155 - loss: 0.2390

Epoch 22/30
1447/1447  8s 6ms/step - accuracy: 0.9163 - loss: 0.2337

Epoch 23/30
1447/1447  8s 6ms/step - accuracy: 0.9196 - loss: 0.2337

```

ss: 0.2273
Epoch 24/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9172 - lo
ss: 0.2332
Epoch 25/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9190 - lo
ss: 0.2300
Epoch 26/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9221 - lo
ss: 0.2236
Epoch 27/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9216 - lo
ss: 0.2235
Epoch 28/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9233 - lo
ss: 0.2215
Epoch 29/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9233 - lo
ss: 0.2191
Epoch 30/30
1447/1447 ————— 8s 6ms/step - accuracy: 0.9251 - lo
ss: 0.2148

```

Out[192]: <keras.src.callbacks.history.History at 0x19e03497f70>

In [193]: *# Testing the data*
model.evaluate(X_test, y_test)

```

582/582 ————— 1s 2ms/step - accuracy: 0.8443 - los
s: 0.5030

```

Out[193]: [0.5132427215576172, 0.8417571783065796]

In [194]: *# Saving the model*
model.save('Loan_Approval.h5')

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

In [195]: *# Loading the trained model*
from keras.models import load_model
model_load = load_model('Loan_Approval.h5')

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

In [199]: *# Predicting the values on Test data set*
y_pred = (model_load.predict(X_test)>0.5).astype('int32')

```

582/582 ————— 1s 2ms/step

```

In [204]: y_pred

Out[204]: array([[1],
[1],
[0],
...,
[1],
[0],
[1]])

```

In [201]: # Testing the model on New Input Data
import joblib

new_data = {'Income':500000, 'Age':30, 'Experience':5, 'Married/Single'

# Converting to pandas DataFrame
df_1 = pd.DataFrame([new_data])

# Concatenation of the new data to the original dataset
new_df = pd.concat([X_old, df_1], ignore_index=True)

# Converting the categorical variables to dummy/indicator variables
new_df = pd.get_dummies(new_df, drop_first=True)

# Extracting the last row(new data) for prediction
new_data_1=new_df.tail(1)
new_data_1

# Loading the scaler
scaling = joblib.load('feature_scaling.pkl')
new_data_1_scaled = scaling.transform(new_data_1)


# Loading the model
model_load = load_model('Loan_Approval.h5')

# Making a prediction
output = (model_load.predict(new_data_1_scaled)>0.5).astype('int32')

# Interpreting the final result
if output == 1:
    print("The Client is High Risk")
else:
    print("The Client is Low Risk")

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

1/1  0s 57ms/step
The Client is High Risk

```

In [205]: from sklearn.metrics import accuracy_score, classification_report, c

print(classification_report(y_pred, y_test))

```

	precision	recall	f1-score	support
0	0.81	0.86	0.84	8747
1	0.87	0.83	0.85	9851
accuracy			0.84	18598
macro avg	0.84	0.84	0.84	18598
weighted avg	0.84	0.84	0.84	18598

In []: